

## Q1 a)


```
In [21]: import pandas as pd
          from pandas import Series
          import numpy as np
          import matplotlib.pyplot as plt
          measurement = pd.read_csv('Measurement_Q1.csv', names = ["time", "measure
          measurement
```

Out[21]:

	time	measurement
<b>0</b>	1	1.84
<b>1</b>	2	3.93
<b>2</b>	3	4.00
<b>3</b>	4	5.42
<b>4</b>	5	6.89
...	...	...
<b>145</b>	146	3.65
<b>146</b>	147	2.82
<b>147</b>	148	2.81
<b>148</b>	149	1.64
<b>149</b>	150	0.00

150 rows × 2 columns

```
In [2]: x = measurement.time  
y = measurement.measurement  
plt.plot (x, y)  
plt.xlabel ('Time')  
plt.ylabel ('Measurement')  
plt.show()
```



```
In [3]: import statsmodels.api as sm
```

```
In [4]: measurement_model = sm.tsa.arima.ARIMA(measurement.measurement, order=(0,
```

```
In [5]: output = measurement_model.fit()
```

```
In [6]: output.summary()
```

Out[6]:

#### SARIMAX Results

**Dep. Variable:** measurement **No. Observations:** 150

**Model:** ARIMA(0, 1, 1) **Log Likelihood** -202.609

**Date:** Thu, 13 Oct 2022 **AIC** 409.217

**Time:** 12:54:06 **BIC** 415.225

**Sample:** 0 **HQIC** 411.658

- 150

**Covariance Type:** opg

	coef	std err	z	P> z	[0.025	0.975]
ma.L1	0.7531	0.060	12.573	0.000	0.636	0.871
sigma2	0.8834	0.109	8.080	0.000	0.669	1.098
Ljung-Box (L1) (Q):	0.26	Jarque-Bera (JB):		1.10		
Prob(Q):	0.61	Prob(JB):		0.58		
Heteroskedasticity (H):	1.00	Skew:		-0.21		
Prob(H) (two-sided):	0.99	Kurtosis:		2.98		


Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

## Q1b)

```
In [7]: measurement2=measurement.diff()  
plt.plot(measurement2)
```

```
Out[7]: [<matplotlib.lines.Line2D at 0x7f9f50138df0>,  
<matplotlib.lines.Line2D at 0x7f9f50138d90>]
```



## Q1 c)

```
In [8]: measurement_diff = measurement2.iloc[1:]
measurement_model2 = sm.tsa.arima.ARIMA(measurement_diff.measurement, ord
```

```
In [9]: measurement_model_fit2 = measurement_model2.fit()
```

```
In [10]: measurement_model_fit2.summary()
```

```
Out[10]:
```

SARIMAX Results			
<b>Dep. Variable:</b>	measurement	<b>No. Observations:</b>	149
<b>Model:</b>	ARIMA(0, 0, 1)	<b>Log Likelihood</b>	-202.607
<b>Date:</b>	Thu, 13 Oct 2022	<b>AIC</b>	411.215
<b>Time:</b>	12:54:06	<b>BIC</b>	420.227
<b>Sample:</b>	0	<b>HQIC</b>	414.876
	- 149		

Covariance Type: opg

	coef	std err	z	P> z	[0.025	0.975]
<b>const</b>	-0.0064	0.137	-0.047	0.962	-0.274	0.261
<b>ma.L1</b>	0.7531	0.061	12.392	0.000	0.634	0.872
<b>sigma2</b>	0.8834	0.110	8.052	0.000	0.668	1.098

Ljung-Box (L1) (Q): 0.26 Jarque-Bera (JB): 1.10

Prob(Q): 0.61 Prob(JB): 0.58

Heteroskedasticity (H): 0.99 Skew: -0.21

Prob(H) (two-sided): 0.98 Kurtosis: 2.98

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).



## Q1 d)

After implementing all parts above it can be observed that both methods give basically the same result because they are doing the same thing. a) part is more of a direct approach whereas, b) is a longer process. We can observe one more thing that is AIC which tells the better fit for a model. The lower the AIC better the model will fit. Therefore, a) has lower AIC means it is likely that the model will fit better through this method.


## Q2 a)

```
In [11]: df=pd.read_csv('GlobalAirTemperature.csv')  
df.head()
```

Out[11]:

	Year	Anomaly, C
0	1880	-0.11
1	1881	-0.13
2	1882	-0.01
3	1883	-0.04
4	1884	-0.42

```
In [12]: x = df['Year']
y = df['Anomaly, C']
plt.plot(x, y)
plt.xlabel('Year')
plt.ylabel('Anomaly')
plt.show()
```



```
In [13]: anomaly_model = sm.tsa.arima.ARIMA(df['Anomaly, C'], order=(0,1,1))
```

```
In [14]: output2=anomaly_model.fit()
```

```
In [15]: output2
```

Out[15]: <statsmodels.tsa.arima.model.ARIMAResultsWrapper at 0x7f9f70ce7760>

In [16]: output2.summary()

Out[16]:

SARIMAX Results

<b>Dep. Variable:</b>	Anomaly, C	<b>No. Observations:</b>	125
<b>Model:</b>	ARIMA(0, 1, 1)	<b>Log Likelihood</b>	72.910
<b>Date:</b>	Thu, 13 Oct 2022	<b>AIC</b>	-141.821
<b>Time:</b>	12:54:06	<b>BIC</b>	-136.180
<b>Sample:</b>	0	<b>HQIC</b>	-139.530
	- 125		
<b>Covariance Type:</b>	opg		

	coef	std err	z	P> z	[0.025	0.975]
<b>ma.L1</b>	-0.6640	0.072	-9.243	0.000	-0.805	-0.523
<b>sigma2</b>	0.0180	0.002	8.182	0.000	0.014	0.022

<b>Ljung-Box (L1) (Q):</b>	2.36	<b>Jarque-Bera (JB):</b>	2.20
<b>Prob(Q):</b>	0.12	<b>Prob(JB):</b>	0.33
<b>Heteroskedasticity (H):</b>	0.87	<b>Skew:</b>	-0.29
<b>Prob(H) (two-sided):</b>	0.65	<b>Kurtosis:</b>	3.32


Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [17]: anomaly_predict = output2.predict (1, len(y), typ = 'levels').rename("Pre
```

```
In [18]: plt.plot (anomaly_predict)
```

```
Out[18]: [<matplotlib.lines.Line2D at 0x7f9f61ef4ac0>]
```



```
In [19]: from sklearn.metrics import mean_squared_error  
mse1 = mean_squared_error (y, anomaly_predict)  
mse1
```

```
Out[19]: 0.007848867110038906
```

```
In [22]: sse=np.sum((y-anomaly_predict)**2)
         sse
```

```
Out[22]: 2.2339840576983625
```

## Q2 b)

```
In [23]: anomaly_model2 = sm.tsa.arima.ARIMA(df['Anomaly, C'], order=(0,1,2))
```

```
In [24]: output_ima2 = anomaly_model2.fit()
```

```
In [25]: output_ima2.summary()
```

```
Out[25]:
```

SARIMAX Results			
<b>Dep. Variable:</b>	Anomaly, C	<b>No. Observations:</b>	125
<b>Model:</b>	ARIMA(0, 1, 2)	<b>Log Likelihood</b>	76.508
<b>Date:</b>	Thu, 13 Oct 2022	<b>AIC</b>	-147.016

<b>Time:</b>	12:55:29				<b>BIC</b>	-138.555
<b>Sample:</b>	0				<b>HQIC</b>	-143.579
- 125						
<b>Covariance Type:</b>	opg					
	<b>coef</b>	<b>std err</b>	<b>z</b>	<b>P&gt; z </b>	<b>[0.025</b>	<b>0.975]</b>
<b>ma.L1</b>	-0.4676	0.099	-4.736	0.000	-0.661	-0.274
<b>ma.L2</b>	-0.2296	0.090	-2.565	0.010	-0.405	-0.054
<b>sigma2</b>	0.0170	0.002	7.631	0.000	0.013	0.021
<b>Ljung-Box (L1) (Q):</b>	0.13	<b>Jarque-Bera (JB):</b>	1.18			
<b>Prob(Q):</b>	0.72	<b>Prob(JB):</b>	0.55			
<b>Heteroskedasticity (H):</b>	1.04	<b>Skew:</b>	-0.23			
<b>Prob(H) (two-sided):</b>	0.91	<b>Kurtosis:</b>	3.12			




Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [26]: anomaly_predict2 = output_ima2.predict (1, len(y), typ = 'levels').rename
```

```
In [27]: plt.plot(anomaly_predict2)
```

```
Out[27]: [<matplotlib.lines.Line2D at 0x7f9f710b0f40>]
```



```
In [28]: mse2 = mean_squared_error (y, anomaly_predict2)
mse2
```

```
Out[28]: 0.004493855486524853
```

```
In [29]: sse=np.sum((y-anomaly_predict2)**2)
         sse
```

```
Out[29]: 2.1120405362089145
```

## Q2 c)

After implementing IMA(1,1) and IMA(1,2) it was observed that sse for order 2 was less, therefore, the model was more close to original values than order 1. Also AIC value was less for order 2 and from above explanation we know that lower AIC is better. Ultimately, we can conclude that for this data IMA(1,2) would be a better fit.

## Q3 a)


```
In [30]: dfmeasure = pd.read_csv('Measurement_Q3.csv')
         dfmeasure.head()
```

Out[30]:


	Year	Measurement
0	1950	2.429415
1	1951	2.363364
2	1952	2.374305
3	1953	2.295520
4	1954	2.329716

```
In [31]: x = dfmeasure.Year  
         y = dfmeasure.Measurement  
         plt.plot(x,y)
```

Out[31]: []



```
In [50]: from statsmodels.graphics.tsaplots import plot_acf
dfmeasure2 = dfmeasure[['Year', 'Measurement']].set_index(['Year'])
plot_acf(dfmeasure2, lags=50);
```



```
In [51]: dfmeasure2
```

```
Out[51]:      Measurement
```

```
      Year
```

```
1950      2.429415
```

<b>1951</b>	2.363364
<b>1952</b>	2.374305
<b>1953</b>	2.295520
<b>1954</b>	2.329716
<b>1955</b>	2.233017
<b>1956</b>	2.378179
<b>1957</b>	2.322671
<b>1958</b>	2.416556
<b>1959</b>	2.498199
<b>1960</b>	2.579453
<b>1961</b>	2.580840
<b>1962</b>	2.629293
<b>1963</b>	2.581853
<b>1964</b>	2.720940

<b>1965</b>	2.844774
<b>1966</b>	3.144862
<b>1967</b>	3.433044
<b>1968</b>	3.358418
<b>1969</b>	3.462620
<b>1970</b>	3.647342
<b>1971</b>	3.991080
<b>1972</b>	3.925702
<b>1973</b>	4.490962
<b>1974</b>	4.491541
<b>1975</b>	4.396567
<b>1976</b>	3.984491
<b>1977</b>	4.115111
<b>1978</b>	4.018538




<b>1979</b>	4.201107
<b>1980</b>	4.367459
<b>1981</b>	4.228103
<b>1982</b>	4.145889
<b>1983</b>	3.786691
<b>1984</b>	3.827373
<b>1985</b>	3.820376
<b>1986</b>	4.025134
<b>1987</b>	4.070130
<b>1988</b>	4.105920
<b>1989</b>	4.039027
<b>1990</b>	4.113978
<b>1991</b>	4.410670
<b>1992</b>	4.110586

<b>1993</b>	4.357700
<b>1994</b>	3.959040
<b>1995</b>	3.882907
<b>1996</b>	3.524803
<b>1997</b>	3.249564
<b>1998</b>	3.139884
<b>1999</b>	3.034263
<b>2000</b>	2.805041
<b>2001</b>	3.304467
<b>2002</b>	2.797697
<b>2003</b>	2.662227
<b>2004</b>	2.589383

```
In [36]: from statsmodels.graphics.tsaplots import plot_pacf  
plot_pacf(dfmeasure2);
```

```
/Users/dhavalgarg/opt/anaconda3/lib/python3.9/site-packages/statsmodels/g  
raphics/tsaplots.py:348: FutureWarning: The default method 'yw' can produ  
ce PACF values outside of the [-1,1] interval. After 0.13, the default wi  
ll change to unadjusted Yule-Walker ('ywm'). You can use this method now b  
y setting method='ywm'.
```

```
warnings.warn(
```



## Q3 b)

```
In [52]: me = diff(dfmeasure2)
```

```
In [54]: me
```

```
Out[54]:
```

	Measurement
1951	-0.066051
1952	0.010941
1953	-0.078785
1954	0.034196
1955	-0.096699
1956	0.145162
1957	-0.055508

	Measurement
1951	-0.066051
1952	0.010941
1953	-0.078785
1954	0.034196
1955	-0.096699
1956	0.145162
1957	-0.055508

<b>1958</b>	0.093885
<b>1959</b>	0.081643
<b>1960</b>	0.081254
<b>1961</b>	0.001387
<b>1962</b>	0.048453
<b>1963</b>	-0.047440
<b>1964</b>	0.139087
<b>1965</b>	0.123834
<b>1966</b>	0.300088
<b>1967</b>	0.288182
<b>1968</b>	-0.074626
<b>1969</b>	0.104202
<b>1970</b>	0.184722
<b>1971</b>	0.343738

<b>1972</b>	-0.065378
<b>1973</b>	0.565260
<b>1974</b>	0.000579
<b>1975</b>	-0.094974
<b>1976</b>	-0.412076
<b>1977</b>	0.130620
<b>1978</b>	-0.096573
<b>1979</b>	0.182569
<b>1980</b>	0.166352
<b>1981</b>	-0.139356
<b>1982</b>	-0.082214
<b>1983</b>	-0.359198
<b>1984</b>	0.040682
<b>1985</b>	-0.006997

<b>1986</b>	0.204758
<b>1987</b>	0.044996
<b>1988</b>	0.035790
<b>1989</b>	-0.066893
<b>1990</b>	0.074951
<b>1991</b>	0.296692
<b>1992</b>	-0.300084
<b>1993</b>	0.247114
<b>1994</b>	-0.398660
<b>1995</b>	-0.076133
<b>1996</b>	-0.358104
<b>1997</b>	-0.275239
<b>1998</b>	-0.109680
<b>1999</b>	-0.105621

<b>2000</b>	-0.229222
<b>2001</b>	0.499426
<b>2002</b>	-0.506770
<b>2003</b>	-0.135470
<b>2004</b>	-0.072844


```
In [53]: from statsmodels.tsa.statespace.tools import diff  
# measurement2= dfmeasure.Measurement.diff()
```

```
Out[53]:
```


	<b>Year</b>	<b>Measurement</b>
<b>1</b>	1.0	-0.066051
<b>2</b>	1.0	0.010941
<b>3</b>	1.0	-0.078785
<b>4</b>	1.0	0.034196
<b>5</b>	1.0	-0.096699



```
In [55]: plt.plot (me)
plt.show()
```




```
In [58]: # pd.plotting.autocorrelation_plot(measurement2[1:])
dfmeasure3 = measurement_diff[['time', 'measurement']].set_index(['time'])
plot_acf(me, lags=40);
```



```
In [59]: plot_pacf(me, lags=20);
```

```
/Users/dhavalgarg/opt/anaconda3/lib/python3.9/site-packages/statsmodels/graphics/tsaplots.py:348: FutureWarning: The default method 'yw' can produce PACF values outside of the [-1,1] interval. After 0.13, the default will change to unadjusted Yule-Walker ('ywm'). You can use this method now by setting method='ywm'.
```

```
warnings.warn(
```



## Q3 c)

We cannot tell the order of the model just by looking at the ACF and PACF plots. We can determine that the data will use ARMA model.

One way of finding the order  $p, q$  can be to build models for different values of  $p, q$  and compare the errors as well as AIC values. The lower the AIC value the better the model is.

In [ ]: